

CS 211 : Tues 02/20 (lecture 14)



Prof. Hummel
(he/him)

- Topics: searching, sorting, map

February 2024

Sunday	Monday	Tuesday	Wednesday	Thursday	Friday	Saturday
				1	2	3
4	5	6	7	8	9	10
11	12	13	14	15	16	17
18	19	20	21	22	23	24
25	26	27	28	29		

www.a-printable-calendar.com

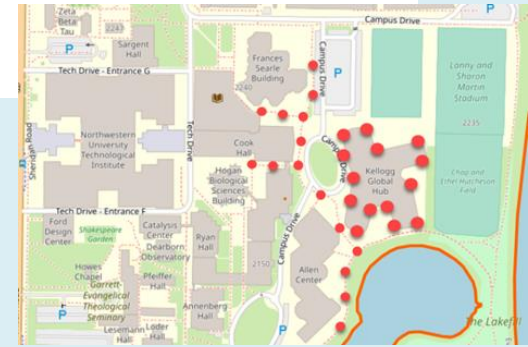
Notes:

- *Lecture slides available on Canvas*
- ***No class Thursday!***
- *HW 06 due tonight (Tuesday)*
- *Project 06 due Friday night (can submit as late as Sunday with late days)*



Northwestern
University

Projects 05 – 08



- Working with Open Street Maps
- Working with EECS computers (aka Linux)

– *Recommendation: open two instances of VS code, editor + terminal*

A screenshot of the Visual Studio Code editor interface. The Explorer pane on the left shows a project named 'proj05' with files like 'a.out', 'building.cpp', 'building.h', 'buildings.cpp', 'buildings.h', 'main.cpp', 'makefile', 'node.cpp', 'node.h', 'nodes-v1.cpp', 'nodes.cpp', and 'nodes.h'. The 'nodes.h' file is selected and its content is displayed in the main editor. The code defines a 'Node' struct with 'id', 'lat', 'lon', and 'isEntrance' fields, and a 'MapNodes' vector. It also includes a 'find' function to search for a node by ID and a 'getNumMapNodes' function to return the count of nodes.

```
1 // nodes.h
2 #ifndef NODES_H
3 #define NODES_H
4
5 #include <iostream>
6 #include <vector>
7 #include <string>
8 #include <map>
9 #include <xml.h>
10
11 // Node struct
12 struct Node {
13     long id;
14     double lat;
15     double lon;
16     bool isEntrance;
17 };
18
19 // Keeps track of all the nodes in the map.
20 class MapNodes {
21 private:
22     vector<Node> MapNodes;
23 public:
24     // readMapNodes
25     //
26     // Given an XML document, reads through the document and
27     // stores all the nodes into the given vector. Each node
28     // is a point on the map, with a unique id along with
29     // (lat, lon) position. Some nodes are entrances to buildings,
30     // which we capture as well.
31     //
32     void readMapNodes(XMLDocument& xmlDoc);
33
34     // find
35     //
36     // Searches the nodes for the one with the matching ID, returning
37     // true if found and false if not. If found, the node's lat, lon,
38     // and isEntrance data are returned via the reference parameters.
39     //
40     bool find(long long id, double& lat, double& lon, bool& isEntrance) const;
41
42     // accessors / getters
43     //
44     int getNumMapNodes() const;
45 };
46
47 #endif
```

A screenshot of the Visual Studio Code terminal window. The terminal shows the output of a program run in a Linux environment. The user has entered 'make run' and the program has executed. It prompts for a map filename ('nu.osm') and displays the number of nodes (15070) and buildings (103). It then prompts for a building name and displays a list of buildings with their coordinates and whether they are entrances. The user has entered 'University Hall' and the program has displayed the corresponding building ID (33908928) and coordinates (42.0518, -87.6758).

```
hummel@moore$ make run
./a.out
** NU open street map **

Enter map filename>
nu.osm
# of nodes: 15070
# of buildings: 103

Enter building name (partial or complete), or * to list, or $ to end>
University Hall
Address: 1897 Sheridan Road
Building ID: 33908928
Nodes:
388499432: (42.0518, -87.6758)
4774714375: (42.0518, -87.6758)
2241369266: (42.0518, -87.6758)
2241369264: (42.0518, -87.6759)
2241227052: (42.0519, -87.6758), is entrance
4774714382: (42.0519, -87.6758)
4774714383: (42.0519, -87.6758)
388499433: (42.052, -87.6758)
388499434: (42.0521, -87.676)
1766764521: (42.0519, -87.6761), is entrance
4774714381: (42.0519, -87.6761)
4774714380: (42.0519, -87.6761)
388499436: (42.0518, -87.6762)
4774714372: (42.0518, -87.676)
2241226778: (42.0518, -87.676)
2241227054: (42.0518, -87.676), is entrance
2241226814: (42.0517, -87.676)
4774714373: (42.0518, -87.676)
4774714374: (42.0517, -87.6759)
4774714376: (42.0517, -87.6759)
4774714377: (42.0517, -87.6758)
4774714379: (42.0517, -87.6758)
4774714378: (42.0517, -87.6758)
388499432: (42.0518, -87.6758)

Enter building name (partial or complete), or * to list, or $ to end>
$
** Done **
```

Nodes

id
lat
lon
entrance

- For "nu.osm", there are 15,000 nodes (positions) in the Nodes vector:



- University Hall has 24 "node refs" outlining the perimeter of the building:

```
Enter building name (partial or complete), or * to list, or $ to end>
University Hall
University Hall
Address: 1897 Sheridan Road
Building ID: 33908928
Nodes: 24
388499432: (42.0518, -87.6758)
4774714375: (42.0518, -87.6758)
2241369266: (42.0518, -87.6758)
2241369264: (42.0518, -87.6759)
2241227052: (42.0519, -87.6758), is entrance
4774714382: (42.0519, -87.6758)
4774714383: (42.0519, -87.6758)
388499433: (42.052, -87.6758)
388499434: (42.0521, -87.676)
1766764521: (42.0519, -87.6761), is entrance
4774714381: (42.0519, -87.6761)
4774714380: (42.0519, -87.6761)
388499436: (42.0518, -87.6762)
4774714372: (42.0518, -87.676)
2241226778: (42.0518, -87.676)
2241227054: (42.0518, -87.676), is entrance
2241226814: (42.0517, -87.676)
4774714373: (42.0518, -87.676)
4774714374: (42.0517, -87.6759)
4774714376: (42.0517, -87.6759)
4774714377: (42.0517, -87.6758)
4774714379: (42.0517, -87.6758)
4774714378: (42.0517, -87.6758)
388499432: (42.0518, -87.6758)
```

Question

id
lat
lon
entrance



- Assume linear search of the Nodes vector, which contains 15,000 elements.
- When outputting "University Hall", the program has to lookup these 24 references to obtain the lat/lon of each node. If the cost of accessing a node in the vector is \$1, how much does it cost (on average) to lookup these 24 nodes?

```
Enter building name (partial or complete), or * to list, or $ to end>
University Hall
University Hall
Address: 1897 Sheridan Road
Building ID: 33908928
Nodes: 24
388499432: (42.0518, -87.6758)
4774714375: (42.0518, -87.6758)
2241369266: (42.0518, -87.6758)
2241369264: (42.0518, -87.6759)
2241227052: (42.0519, -87.6758), is entrance
4774714382: (42.0519, -87.6758)
4774714383: (42.0519, -87.6758)
388499433: (42.052, -87.6758)
388499434: (42.0521, -87.676)
1766764521: (42.0519, -87.6761), is entrance
4774714381: (42.0519, -87.6761)
4774714380: (42.0519, -87.6761)
388499436: (42.0518, -87.6762)
4774714372: (42.0518, -87.676)
2241226778: (42.0518, -87.676)
2241227054: (42.0518, -87.676), is entrance
2241226814: (42.0517, -87.676)
4774714373: (42.0518, -87.676)
4774714374: (42.0517, -87.6759)
4774714376: (42.0517, -87.6759)
4774714377: (42.0517, -87.6758)
4774714379: (42.0517, -87.6758)
4774714378: (42.0517, -87.6758)
388499432: (42.0518, -87.6758)
```

- A) \$300
- B) \$359,700
- C) \$180,000
- D) \$24,000
- E) Over \$1,000,000

Demo: linear search

- Let's confirm the result...


```
** NU open street map **

Enter map filename>
nu.osm
# of nodes: 15070
# of buildings: 103

Enter building name (partial or complete), or * to list, or $ to end>
University Hall
University Hall
Address: 1897 Sheridan Road
Building ID: 33908928
Nodes:
388499432: (42.0518, -87.6758)
4774714375: (42.0518, -87.6758)
2241369266: (42.0518, -87.6758)
2241369264: (42.0518, -87.6759)
2241227052: (42.0519, -87.6758), is entrance
4774714382: (42.0519, -87.6758)
4774714383: (42.0519, -87.6758)
388499433: (42.052, -87.6758)
388499434: (42.0521, -87.676)
1766764521: (42.0519, -87.6761), is entrance
4774714381: (42.0519, -87.6761)
4774714380: (42.0519, -87.6761)
388499436: (42.0518, -87.6762)
4774714372: (42.0518, -87.676)
2241226778: (42.0518, -87.676)
2241227054: (42.0518, -87.676), is entrance
2241226814: (42.0517, -87.676)
4774714373: (42.0518, -87.676)
4774714374: (42.0517, -87.6759)
4774714376: (42.0517, -87.6759)
4774714377: (42.0517, -87.6758)
4774714379: (42.0517, -87.6758)
4774714378: (42.0517, -87.6758)
388499432: (42.0518, -87.6758)


Enter building name (partial or complete), or * to list, or $ to end>
$

** Done **
# of calls to getID(): 120146
# of Nodes created: 15070
# of Nodes copied: 151599
```



Discussion

- **Linear search:**



- **Costs?**

- *Best (lowest) possible cost: $1 + 2 + \dots + 24 = n(n+1)/2 = \300*
- *Worst (highest) possible cost (none found): $24 * 15000 = \$360,000$*
- *On average, you expect the cost to be: $24 * 7,500 = \$180,000$*

- **We say the linear search algorithm has a time complexity "on the order of N"**
 - *as N increases, the time increases*
 - *Written $O(N)$*

Is there a faster way?

- **Yes!**
- **Binary search...**

Demo: binary search

- How much faster?

```
** NU open street map **

Enter map filename>
nu.osm
# of nodes: 15070
# of buildings: 103

Enter building name (partial or complete), or * to list, or $ to end>
University Hall
University Hall
Address: 1897 Sheridan Road
Building ID: 33908928
Nodes:
  388499432: (42.0518, -87.6758)
  4774714375: (42.0518, -87.6758)
  2241369266: (42.0518, -87.6758)
  2241369264: (42.0518, -87.6759)
  2241227052: (42.0519, -87.6758), is entrance
  4774714382: (42.0519, -87.6758)
  4774714383: (42.0519, -87.6758)
  388499433: (42.052, -87.6758)
  388499434: (42.0521, -87.676)
  1766764521: (42.0519, -87.6761), is entrance
  4774714381: (42.0519, -87.6761)
  4774714380: (42.0519, -87.6761)
  388499436: (42.0518, -87.6762)
  4774714372: (42.0518, -87.676)
  2241226778: (42.0518, -87.676)
  2241227054: (42.0518, -87.676), is entrance
  2241226814: (42.0517, -87.676)
  4774714373: (42.0518, -87.676)
  4774714374: (42.0517, -87.6759)
  4774714376: (42.0517, -87.6759)
  4774714377: (42.0517, -87.6758)
  4774714379: (42.0517, -87.6758)
  4774714378: (42.0517, -87.6758)
  388499432: (42.0518, -87.6758)

Enter building name (partial or complete), or * to list, or $ to end>
$

** Done **
# of calls to getID() 316
# of Nodes created: 15070
# of Nodes copied: 31453
```

Binary search

2	4	9	14	18	22	36	54	71	88	101
---	---	---	----	----	----	----	----	----	----	-----

- **Jump to the middle**
- **Compare --- if == stop, if < go left, if > go right**
- **Repeat**

- **Binary search is a divide-and-conquer algorithm, dividing the search space in half each time...**
- **We say the binary search algorithm has a time complexity "on the order of $\log_2 N$ "**
 - *as N increases, the time increases much more slowly*
 - *Written $O(\lg N)$*

Interview question


- There are 65,000 elements in a vector, in random order
- You need to perform 4 searches of this data
- What is the most efficient approach?

- A) Use linear search
- B) Use binary search
- C) Sort then use linear search
- D) Sort then use binary search

Binary search pre-condition: **sorted order**

- To use binary search, data must be sorted!


18	22	9	4	88	36	14	101	2	54	71
----	----	---	---	----	----	----	-----	---	----	----



2	4	9	14	18	22	36	54	71	88	101
---	---	---	----	----	----	----	----	----	----	-----

Sorting in C++

- Built-in algorithm
- Uses lambda function to sort objects




```
#include <algorithm>
```

```
std::sort(this->MapNodes.begin(), this->MapNodes.end(),
    [](const Node& node1, const Node& node2)
    {
        if (node1.getID() < node2.getID()) // keep this order
            return true;
        else // swap them
            return false;
    }
);
```

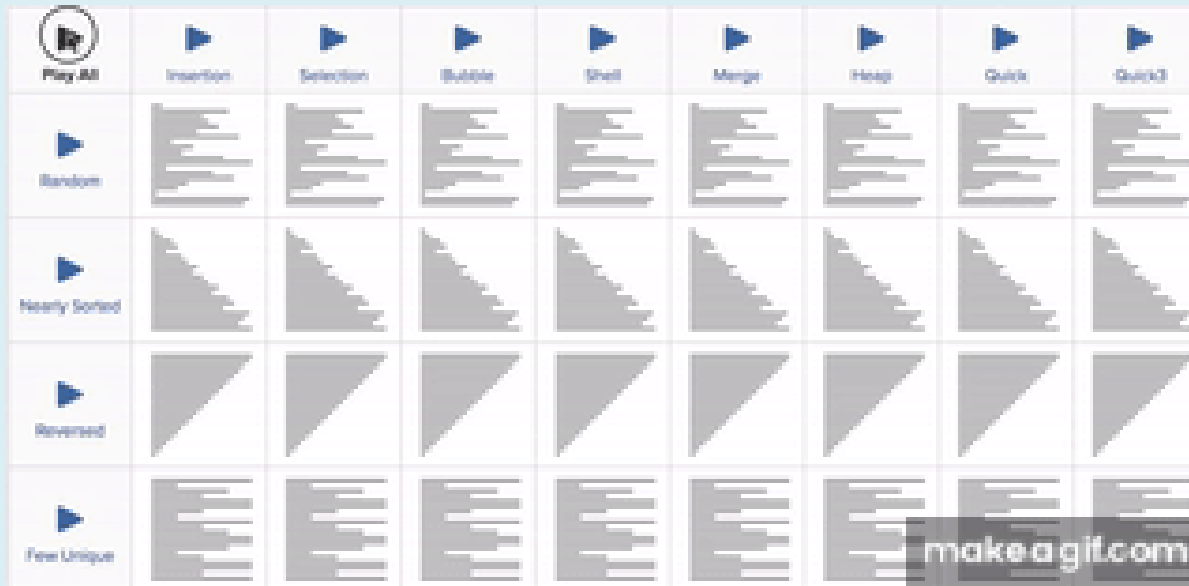
Demo: sorting

- What's the cost of sorting the nodes?

```
** NU open street map **  
  
Enter map filename>  
nu.osm  
# of nodes: 15070  
# of buildings: 103  
  
Enter building name (partial or complete), or * to list, or $ to end>  
$  
  
** Done **  
# of calls to getID(): 500584  
# of Nodes created: 15070  
# of Nodes copied: 33497
```




Sorting Algorithms



Sorting visualizations:

- <http://www.sorting-algorithms.com>
- <http://www.cs.usfca.edu/~galles/visualization/ComparisonSort.html>
- “15 sorts in 6 minutes” video on YouTube:
<https://www.youtube.com/watch?v=kPRA0W1kECg>

Time complexities



Does time complexity really matter?

Example:

– Suppose $N = 1,000,000$ (1MB)


	Time Complexity	# of steps	Example Algorithm
Algorithm A	$O(1)$		
Algorithm B	$O(\lg N)$		
Algorithm C	$O(N)$		
Algorithm D	$O(N \lg N)$		
Algorithm E	$O(N^2)$		



Searching


Searching

linear search



$O(n)$ per search


sort + binary search



$O(n \lg n)$ sort
 $O(\lg n)$ per search

$O(n)$ insert/delete

search trees




$O(n \lg n)$ to build
 $O(\lg n)$ per search

$O(\lg n)$ insert/delete

Question

Is this a valid search tree?



- A) yes
- B) no, 18 is out of place
- C) no, 19 is out of place
- D) no, 31 is out of place

map

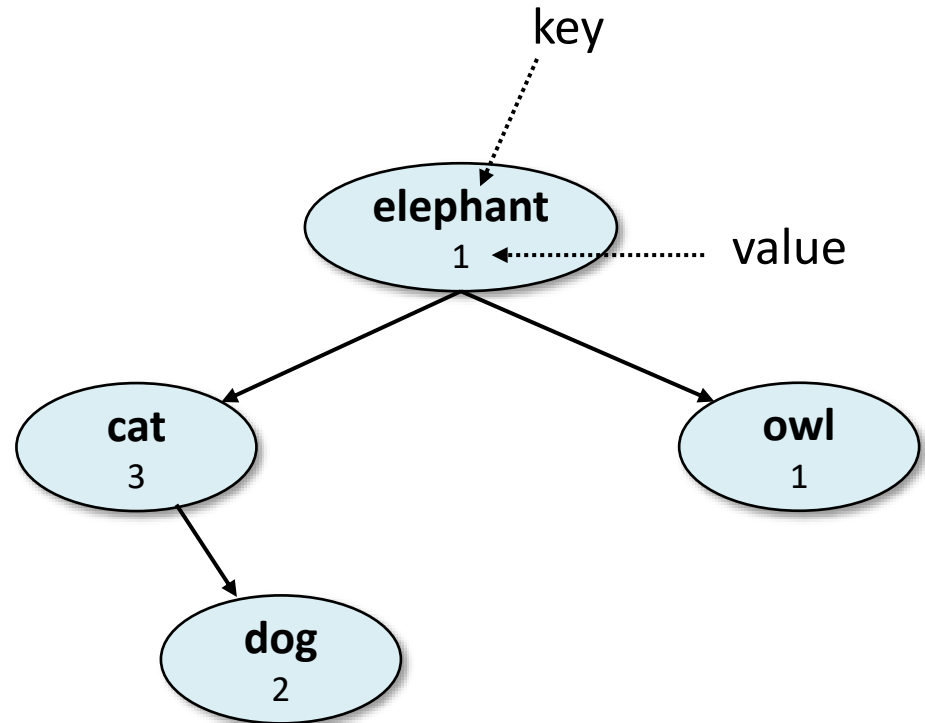

- **map** is C++ abstraction for a search tree

– *"map" key to value*

```
#include <map>
```

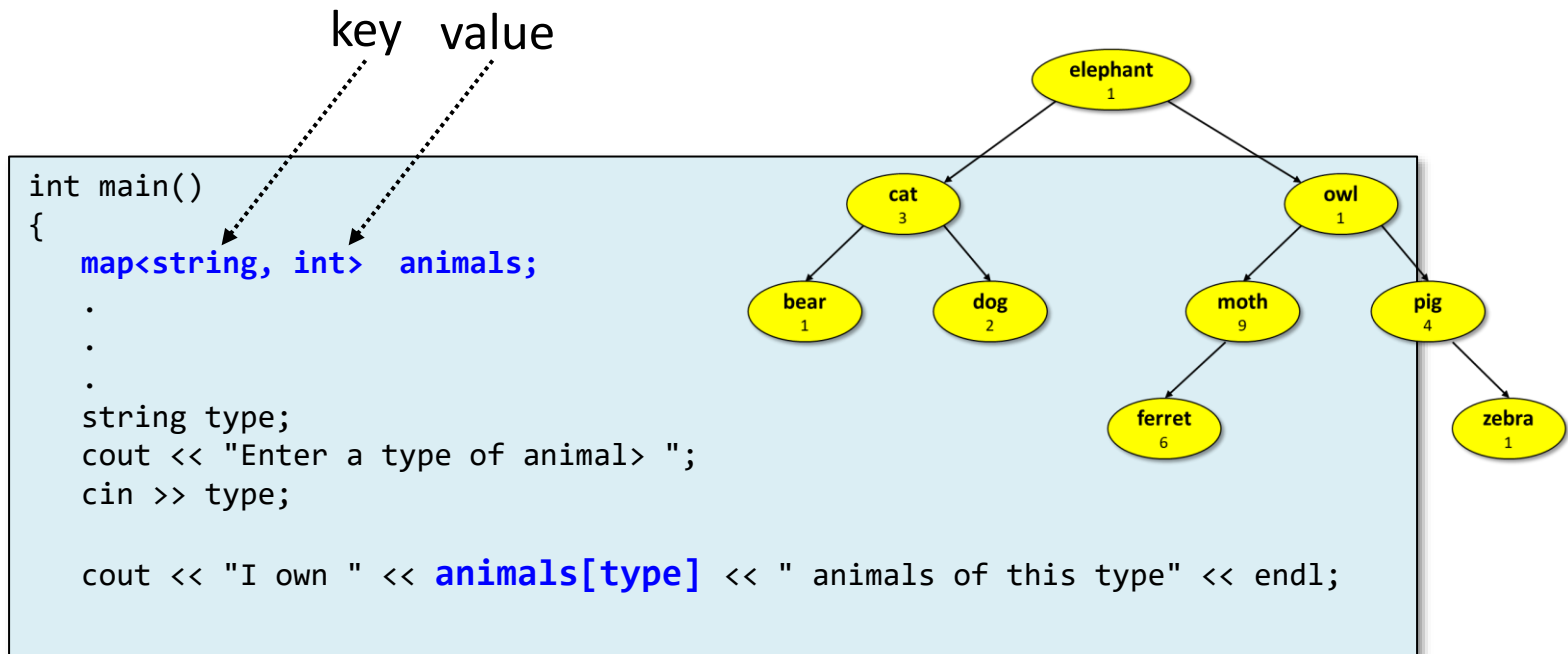

```
int main()
{
    map<string, int> animals;

    animals["elephant"] = 1;
    animals["cat"] = 3;
    animals["owl"] = 1;
    animals["dog"] = 2;
    .
    .
    .
}
```



<key, value> pairs

- Designed for fast $O(\lg N)$ lookup by key



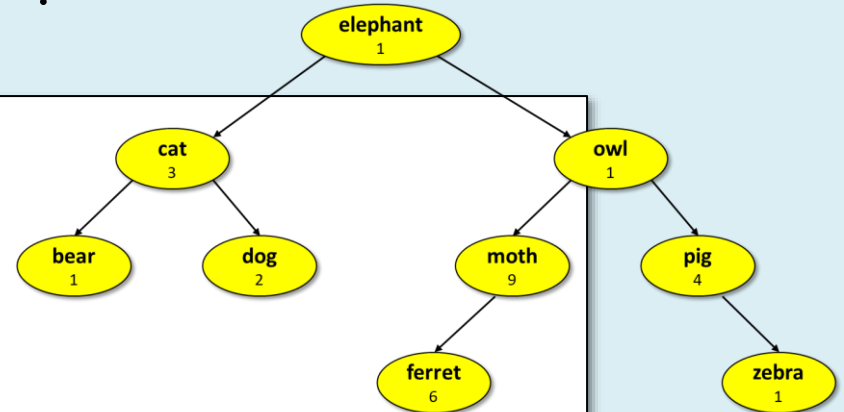
beware of []

- [] performs a search
 - *If not found, inserts key with default value*
- Example:
 - *What if the user enters "newt" ?*

```
int main()
{
    map<string, int>  animals;
    .
    .
    .
    string type;
    cout << "Enter a type of animal> ";
    cin >> type;

    cout << "I own " << animals[type] << " animals of this type" << endl;

    cout << animals.size() << endl;
}
```



prefer searching with `find()`

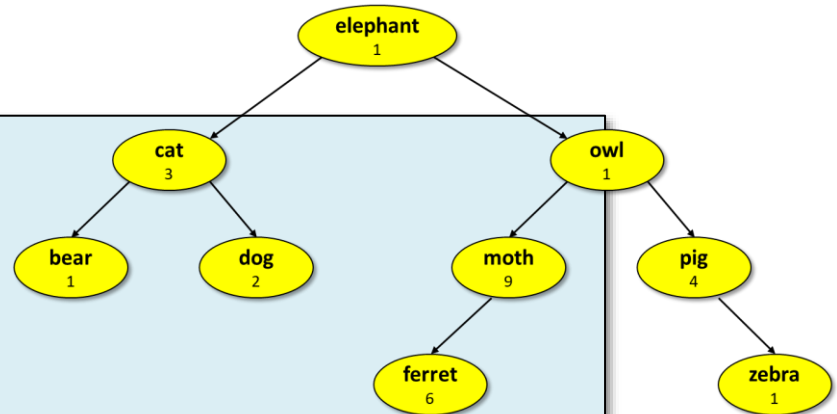
- `find()` will search and return the following:
 - Returns an *iterator* ("pointer") to the `<key, value>` pair if found
 - Returns "end of iteration space" if not found

```
int main()
{
    map<string, int>  animals;
    .
    .
    .
    string type;
    cout << "Enter a type of animal> ";
    cin >> type;

    // cout << "I own " << animals[type] << " animals of this type" << endl;

    auto iter = animals.find(type);

    if (iter == animals.end()) // not found
        cout << "I don't own this type of animal..." << endl;
    else
        cout << "I own " << iter->second << " animals of this type" << endl;
}
```



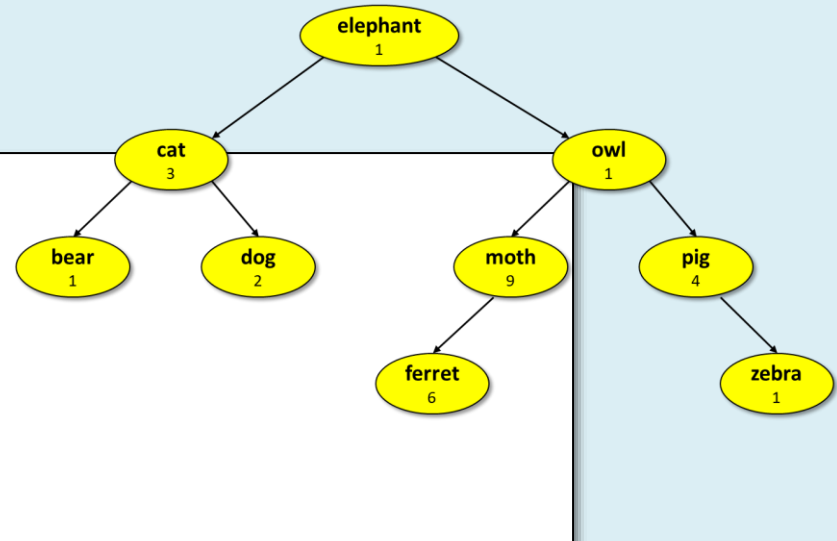
iteration

- Interestingly, you can iterate across the tree using **foreach**
 - *Traverses in key order...*

```
int main()
{
    map<string, int> animals;
    .
    .
    .

    cout << "Animals I own:" << endl;

    for (auto kv_pair : animals)
        cout << kv_pair.first << ": " << kv_pair.second << endl;
}
```



Animals I own:
bear: 1
cat: 3
dog: 2
elephant: 1
ferret: 6
moth: 9
owl: 1
pig: 4
zebra: 1

Humm, I wonder how this works? [project 08?]

map: API summary

Capacity:

empty	Test whether container is empty (public member function)
size	Return container size (public member function)
max_size	Return maximum size (public member function)

Element access:

operator[]	Access element (public member function)
at <small>C++11</small>	Access element (public member function)

} *lookup (beware [] will insert)*

Modifiers:

insert	Insert elements (public member function)
erase	Erase elements (public member function)
swap	Swap content (public member function)
clear	Clear content (public member function)
emplace <small>C++11</small>	Construct and insert element (public member function)
emplace_hint <small>C++11</small>	Construct and insert element with hint (public member function)

Additional ways to insert

Observers:


key_comp	Return key comparison object (public member function)
value_comp	Return value comparison object (public member function)

A better way to lookup...

Operations:

find	Get iterator to element (public member function)
count	Count elements with a specific key (public member function)
lower_bound	Return iterator to lower bound (public member function)
upper_bound	Return iterator to upper bound (public member function)
equal_range	Get range of equal elements (public member function)

*Lookup, insert, and
erase are $O(\lg N)$*



What's due?

HW 06 is due tonight

Project 06 is due Friday night

No class Thursday!

