

CS 211 : Tues 02/27 (lecture 16)



Prof. Hummel
(he/him)

- Topics: efficiency, search trees, set

February 2024

Sunday	Monday	Tuesday	Wednesday	Thursday	Friday	Saturday
				1	2	3
4	5	6	7	8	9	10
11	12	13	14	15	16	17
18	19	20	21	22	23	24
25	26	27	28	29		

www.a-printable-calendar.com

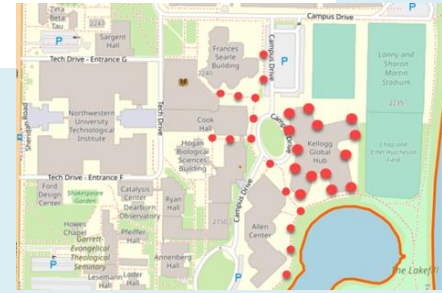
Notes:

- *Lecture slides available on Canvas*
- *HW 07 due tonight (Tuesday)*
- *Project 07 due Friday night (can submit as late as Sunday with late days)*



Northwestern
University

Project 06 --- efficiency



- Imagine working with a real map
 - *Millions of nodes / positions*
 - *Millions of footways / sidewalks / roads*
 - *Millions of buildings / structures*

```
hummel> ./a.out
** NU open street map **

Enter map filename>
ns.osm
# of nodes: 114044
# of buildings: 6916
# of footways: 1486336
Time: 23402ms

Enter building name (partial or complete), or * to list, or $ to end>
University Library
Northwestern University Library
Address: 1970 Campus Drive
Building ID: 175187764
Nodes: 268
Footways that intersect: 3072
Time: 9159ms
```

Correctness first, efficiency second...

Timing in C++

- Add timing code to see what really works...

```
#include <chrono>
```

```
auto start = chrono::high_resolution_clock::now();  
  
.    
.  // computation you want to time  
.    
  
auto stop = chrono::high_resolution_clock::now();  
auto diff = stop - start;  
auto duration = chrono::duration_cast<chrono::milliseconds>(diff);  
  
cout << "Time: " << duration.count() << "ms" << endl;
```

Approach #1

- This takes nearly a minute... And yields duplicate footway IDs... what would you recommend as improvements?

```
void Footways::intersectWithThisNode(long long nodeid, vector<long long>& footwayIDs)
{
    for (Footway F : this->MapFootways) {
        if (F.intersectWithThisNode(nodeid)) {
            footwayIDs.push_back(F.ID);
        }
    }
}
```

```
bool Footway::intersectWithThisNode(long long nodeid)
{
    for (long long id : this->NodeIDs)
        if (id == nodeid)
            return true;

    return false;
}
```

```
void Building::intersectWithFootways(Footways& footways)
{
    vector<long long> footwayIDs;

    auto start = chrono::high_resolution_clock::now();

    for (long long id : this->NodeIDs) { // for each node in this building
        footways.intersectWithThisNode(id, footwayIDs);
    }

    auto stop = chrono::high_resolution_clock::now();
    auto diff = stop - start;
    auto duration = chrono::duration_cast<chrono::milliseconds>(diff);

    cout << "Footways that intersect: " << footwayIDs.size() << endl;
    cout << "Time: " << duration.count() << "ms" << endl;
}
```

```
Enter building name (partial or complete), or * to list, or $ to end>
University Library
Northwestern University Library
Address: 1970 Campus Drive
Building ID: 175187764
Nodes: 268
Footways that intersect: 4096
Time: 56762ms
```

Improvements...

1. For loop should reference footway, not copy
2. Replace vector with set --- eliminates duplicates
3. Store footway nodes in a set for $O(\lg N)$ searching
4. Order matters...

Approach #2

- You can get < 1 second using sets and ordering things properly...

```
bool Footway::intersectWithBuilding(Building& B)
{
    for (long long id : this->NodeIDs)
        if (B.SetIDs.count(id) > 0) // O(lgN) search:
            return true;

    return false;
}
```

```
void Footways::intersectWithBuilding(Building& B)
{
    vector<long long> footwayIDs;


    auto start = chrono::high_resolution_clock::now();

    for (Footway& F : this->MapFootways) { // for each footway...
        if (F.intersectWithBuilding(B)) { // if footway intersects building, store ID:
            footwayIDs.push_back(F.ID);
        }
    }

    auto stop = chrono::high_resolution_clock::now();
    auto diff = stop - start;
    auto duration = chrono::duration_cast<chrono::milliseconds>(diff);

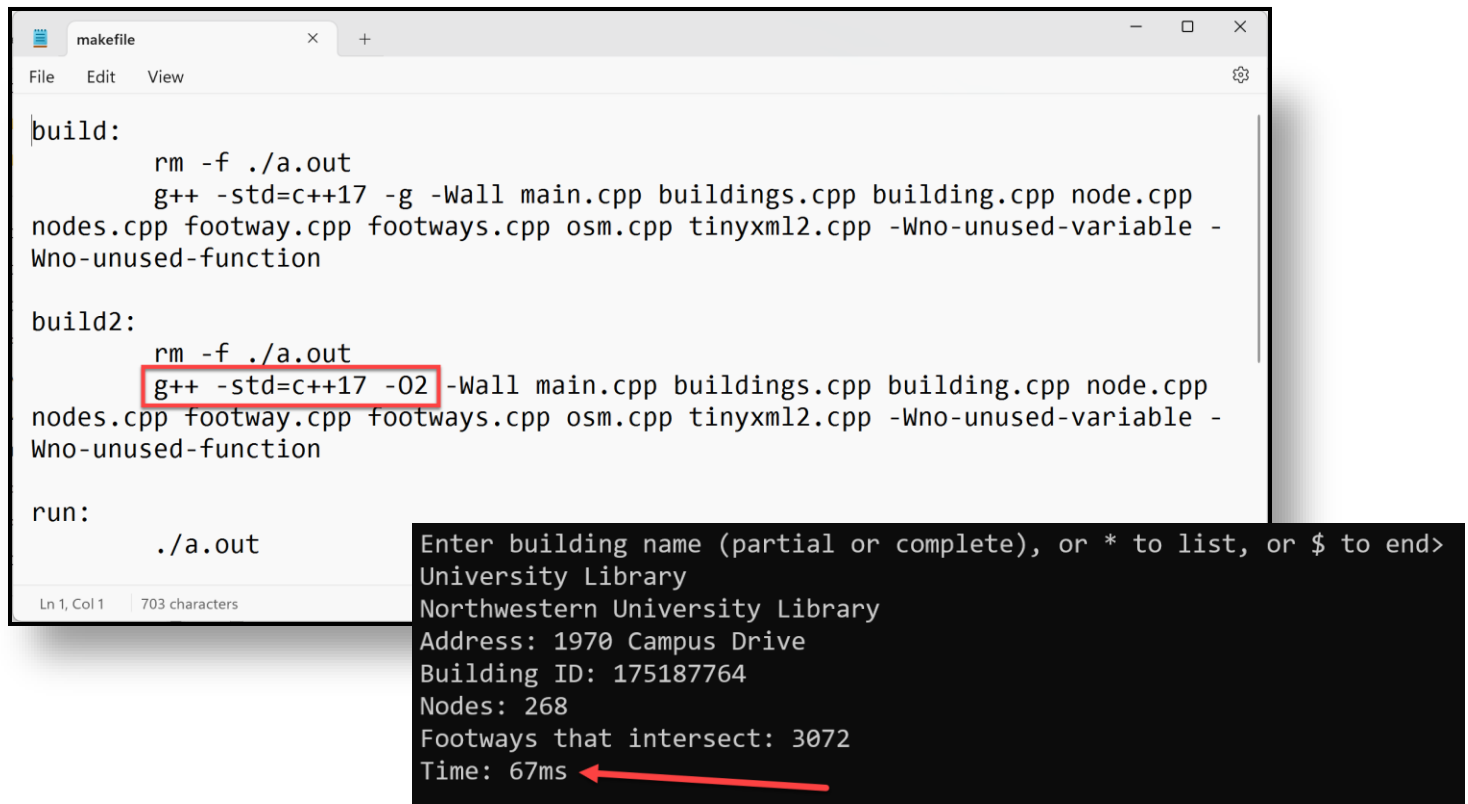
    cout << "Footways that intersect: " << footwayIDs.size() << endl;
    cout << "Time: " << duration.count() << "ms" << endl;
}
```

```
Enter building name (partial or complete), or * to list, or $ to end>
University Library
Northwestern University Library
Address: 1970 Campus Drive
Building ID: 175187764
Nodes: 268
Footways that intersect: 3072
Time: 743ms
```



-02

- Code optimization yields another 10x speed increase



The image shows a code editor window titled 'makefile' with a menu bar (File, Edit, View) and a status bar (Ln 1, Col 1 | 703 characters). The editor contains a Makefile with three targets: 'build:', 'build2:', and 'run:'. The 'build2:' target's command line is highlighted with a red box, showing the addition of the '-O2' optimization flag. Below the editor, a terminal window displays the output of the 'run' target, which includes building name, address, building ID, number of nodes, and the number of intersecting footways. The final line of the terminal output is 'Time: 67ms', with a red arrow pointing to it from the right.

```
makefile
File Edit View

build:
    rm -f ./a.out
    g++ -std=c++17 -g -Wall main.cpp buildings.cpp building.cpp node.cpp
nodes.cpp footway.cpp footways.cpp osm.cpp tinyxml2.cpp -Wno-unused-variable -
Wno-unused-function

build2:
    rm -f ./a.out
    g++ -std=c++17 -O2 -Wall main.cpp buildings.cpp building.cpp node.cpp
nodes.cpp footway.cpp footways.cpp osm.cpp tinyxml2.cpp -Wno-unused-variable -
Wno-unused-function


run:
    ./a.out

Ln 1, Col 1 | 703 characters
```


```
Enter building name (partial or complete), or * to list, or $ to end>
University Library
Northwestern University Library
Address: 1970 Campus Drive
Building ID: 175187764
Nodes: 268
Footways that intersect: 3072
Time: 67ms
```

Approach #3

- "Bounding Box" algorithm
 - For each building & footway, pre-compute "bounding box"
 - Only consider footways that overlap with building




```
Enter building name (partial or complete), or * to list, or $ to end>
University Library
Northwestern University Library
Address: 1970 Campus Drive
Building ID: 175187764
Nodes: 268
Footways that intersect: 3072
Time: 293ms
```

set

- **set** is another C++ abstraction for a search tree
 - *Models mathematical set (no duplicates)*
 - *Efficient operations $\Rightarrow O(\lg N)$*



Set

```
template<typename TKey>
```

```
class set
```

```
{
```

```
private:
```

```
    struct NODE
```

```
    {
```

```
        TKey   Key;
```

```
        NODE*  Left;
```

```
        NODE*  Right;
```

```
    };
```

```
    NODE* Root; // pointer to root node
```

```
    int Size; // # of nodes in tree
```

```
public:
```

```
    // default constructor:
```

```
    set()
```

```
    : Root(nullptr),
```

```
      Size(0)
```

```
    { }
```

```
#include <set>
```

```
int main()
```

```
{
```

```
    set<int> S;
```

```
    S.insert(22);
```

```
    S.insert(11);
```

```
    S.insert(49);
```

```
    .
```


```
    .
```

```
    .
```


Root

Size

9



- Search: You are searching for the key **40**. Trace out the path searched, highlighting the edges traversed & circling the nodes visited.



set::contains(key)


- Searches set for given key

```
int main()
{
    set<int> S;

    S.insert(22);
    S.insert(11);
    S.insert(49);
    .
    .
    .

    int x;
    cout << "Enter an integer> ";
    cin >> x;

    if (S.contains(x))
        cout << x << " is a member" << endl;
}
```



set::contains

```
template<typename TKey>
class set
{
private:
    struct NODE
    {
        TKey  Key;
        NODE* Left;
        NODE* Right;
    };

    NODE* Root; // pointer to
    int   Size; // # of nodes


public:
    // default constructor:
    set()
        : Root(nullptr),
          Size(0)
    { }
};
```

```
bool contains(TKey key)
```


```
{
    NODE* cur = this->Root;

    while (cur != nullptr)
    {
        if (key == cur->Key) { // found it!
            return true;
        }
        else if (key < cur->Key) { // search left:
            cur = cur->Left;
        }
        else { // search right:
            cur = cur->Right;
        }
    }

    // if get here, not found
    return false;
}
```



1) Here's a *mystery* function in the set class. What value is returned for the tree shown?



```
TKey mystery()
{
    NODE* cur = this->Root;

    if (cur == nullptr)
        throw runtime_error("mystery: empty");

    while (cur->Right != nullptr)
    {
        cur = cur->Right;
    }

    return cur->Key;
}
```

A) 49


B) 61

C) 58

D) Throws runtime error


Inserting into a set

- Elements are always inserted at the bottom
 - *as a new “leaf”*



- **Algorithm:**
 - *Start by searching the tree...*
 - *If you find key, then return (don't insert twice)*
 - *Else insert new node where you fell out of the tree...*

- Example: Insert 44



set::insert

```
void insert(TKey key)
```


```
{
```

```
    // STEP 1: search for key, return if
    // found since no duplicates
    while (...)
    { ... }
```

```
    // STEP 2: if not found, insert where
    // we fell out of the tree
    NODE *n = new NODE();
    link in...
```

```
    // STEP 3: update size and return
    this->Size++;
    return;
```

```
}
```



```
template<typename TKey>
class set
{
private:
    struct NODE
    {
        TKey   Key;
        NODE*  Left;
        NODE*  Right;
    };


    NODE*  Root; // pointer to root node
    int    Size; // # of nodes in tree

public:
    // default constructor:
    set()
    : Root(nullptr),
      Size(0)
    { }
```

Step 1

```
void insert(TKey key)
{
    NODE* prev = nullptr;
    NODE* cur = this->Root;

    //
    // 1. Search for key, return if found:
    //
    while (cur != nullptr)
    {
        if (key == cur->Key) { // found
            return;
        }
        else if (key < cur->Key) { // left:
            prev = cur;
            cur = cur->Left;
        }
        else { // right:
            prev = cur;
            cur = cur->Right;
        }
    }
}
```




Step 2

```
void insert(TKey key)
{
    NODE* prev = nullptr;
    NODE* cur = Root;

    .
    . // search...
    .


    //
    // 2. If not found, insert where we
    //    fell out of the tree:
    //
    NODE *n = new NODE();
    n->Key = key;
    n->Left = nullptr;
    n->Right = nullptr;

    << continued on next page >>
}
```



2) The last bit of **step 2** is to link in the new node. Which is the correct code fragment that works in all cases?

Root
Size 9



.
. .
.

```
//  
// 2. If not found, insert where we  
// fell out of the tree:  
//
```

```
NODE *n = new NODE();  
n->Key = key;  
n->Left = nullptr;  
n->Right = nullptr;
```

```
if (key < prev->Key)  
    prev->Left = n;  
else if (key > prev->Key)  
    prev->Right = n;  
else  
    this->Root = n; (A)
```

```
if (prev == nullptr)  
    this->Root = n;  
else if (key < prev->Key)  
    prev->Left = n;  
else  
    prev->Right = n; (B)
```

```
if (cur == nullptr)  
    this->Root = n;  
else if (key < cur->Key)  
    cur->Left = n;  
else  
    cur->Right = n; (C)
```

What's due?

HW 07 is due tonight

Project 07 is due Friday night

